

Assignment 1

All the activation functions and their derivatives

```
import numpy as np
import matplotlib.pyplot as plt

# Define activation functions and their derivatives
def linear(x):
    return x

def linear_derivative(x):
    return np.ones_like(x)

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):
    s = sigmoid(x)
    return s * (1 - s)

def tanh(x):
    return np.tanh(x)

def tanh_derivative(x):
    return 1 - np.tanh(x)**2

def softmax(x):
    e_x = np.exp(x - np.max(x)) # For numerical stability
    return e_x / e_x.sum(axis=0)

def softmax_derivative(x):
    s = softmax(x)
    return s * (1 - s) # Simplified for 1D case, not a full Jacobian

def relu(x):
    return np.maximum(0, x)

def relu_derivative(x):
    return np.where(x > 0, 1, 0)

def leaky_relu(x, alpha=0.01):
    return np.where(x > 0, x, alpha * x)

def leaky_relu_derivative(x, alpha=0.01):
    return np.where(x > 0, 1, alpha)

# Generate input range
x = np.linspace(-5, 5, 500)
```

```

# Prepare for plotting
activations = {
    "Linear": (linear, linear_derivative),
    "Sigmoid": (sigmoid, sigmoid_derivative),
    "Tanh": (tanh, tanh_derivative),
    "Softmax": (softmax, softmax_derivative),
    "ReLU": (relu, relu_derivative),
    "Leaky ReLU": (leaky_relu, leaky_relu_derivative)
}

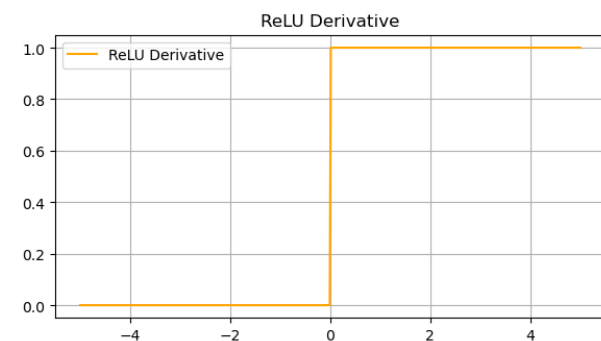

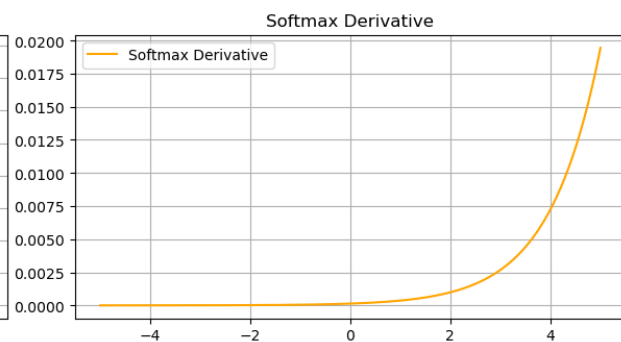





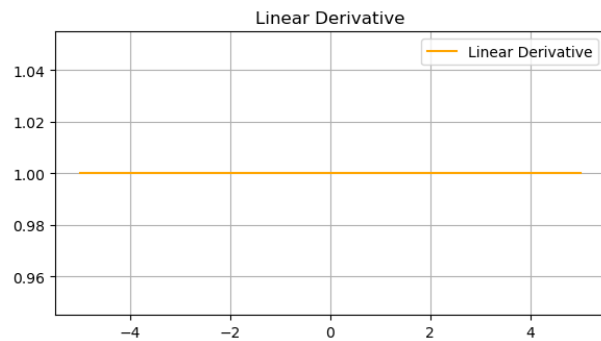

# Plot activation functions and their derivatives
fig, axes = plt.subplots(len(activations), 2, figsize=(12, 20))
for i, (name, (func, deriv)) in enumerate(activations.items()):
    # Compute function and derivative values
    y = func(x)
    dydx = deriv(x)

    # Plot activation function
    axes[i, 0].plot(x, y, label=f"{name} Function")
    axes[i, 0].set_title(f"{name} Activation Function")
    axes[i, 0].legend()
    axes[i, 0].grid(True)

    # Plot derivative
    axes[i, 1].plot(x, dydx, label=f"{name} Derivative",
color='orange')
    axes[i, 1].set_title(f"{name} Derivative")
    axes[i, 1].legend()
    axes[i, 1].grid(True)

plt.tight_layout()
plt.show()

```



```

import pandas as pd
import numpy as np

# Set random seed for reproducibility
np.random.seed(42)

# Number of samples
num_samples = 1000

soil_texture = np.random.uniform(0, 1, num_samples) # Ratio of sand,
silt, and clay
pH = np.random.uniform(4.5, 9.0, num_samples) # Soil pH levels
moisture_content = np.random.uniform(5, 50, num_samples) # Percentage
moisture
organic_matter = np.random.uniform(1, 10, num_samples) # Organic
matter in %
bulk_density = np.random.uniform(1.1, 1.6, num_samples) # Soil bulk
density in g/cm^3

soil_types = []
for i in range(num_samples):
    if soil_texture[i] > 0.7 and pH[i] > 6.5 and moisture_content[i] <
15:
        soil_types.append("Sandy")
    elif organic_matter[i] > 5 and moisture_content[i] > 30:
        soil_types.append("Loamy")
    else:
        soil_types.append("Clayey")

# Create a DataFrame
data = pd.DataFrame({
    "Soil_Texture": soil_texture,
    "pH": pH,
    "Moisture_Content": moisture_content,
    "Organic_Matter": organic_matter,
    "Bulk_Density": bulk_density,
    "Soil_Type": soil_types
})

# Save the dataset to a CSV file
data.to_csv("synthetic_soil_dataset.csv", index=False)

print("Synthetic dataset created and saved as
'synthetic_soil_dataset.csv'.")

Synthetic dataset created and saved as 'synthetic_soil_dataset.csv'.

```

Data Preprocessing:

```

import pandas as pd
import numpy as np
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, LabelEncoder

# Load the synthetic soil dataset
data = pd.read_csv("synthetic_soil_dataset.csv")

# Introduce some missing values for demonstration purposes
data.iloc[10:15, 2] = np.nan # Simulate missing values in
'Moisture_Content'

# 1. Handle Missing Data
imputer = SimpleImputer(strategy="mean") # Replace missing values
with the column mean
data.iloc[:, :-1] = imputer.fit_transform(data.iloc[:, :-1]) # Handle
numerical columns

# 2. Separate Features and Labels
X = data.iloc[:, :-1].values # Features (inputs)
y = data.iloc[:, -1].values # Labels (outputs)

# 3. Encode Categorical Labels (Simple Label Encoding)
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

# 4. Scale the Features (Standardization)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Output Preprocessed Data
print("Features after scaling:\n", X_scaled[:5])
print("Encoded Labels:\n", y_encoded[:5])

# Optional: Save the preprocessed data
preprocessed_data = pd.DataFrame(X_scaled, columns=["Soil_Texture",
"pH", "Moisture_Content", "Organic_Matter", "Bulk_Density"])
preprocessed_data["Soil_Type"] = y_encoded
preprocessed_data.to_csv("preprocessed_soil_dataset.csv", index=False)

print("\nPreprocessed data saved as 'preprocessed_soil_dataset.csv'.")

Features after scaling:
[[-0.39630103 -1.10217857 -0.83490537  0.63673984  0.27171114 ]
 [ 1.57695733  0.11944663 -0.8858098  1.06970686  1.08602506]
 [ 0.82789256  1.25299211  1.39301993 -0.83782241  0.92810149]
 [ 0.37125061  0.77114327 -0.87693542  0.46970805 -1.18676377]
 [-1.14468466  1.02568136 -0.79949617  0.28416991 -1.20298622]]
Encoded Labels:
[0 0 0 0 0]

```

Preprocessed data saved as 'preprocessed_soil_dataset.csv'.

```
dataset = pd.read_csv('preprocessed_soil_dataset.csv')
```

2. Train-Test Split:

```
from sklearn.model_selection import train_test_split

# Perform Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y_encoded, test_size=0.2, random_state=42,
    stratify=y_encoded
)

# Output the shapes of the splits
print(f"Training Features Shape: {X_train.shape}")
print(f"Test Features Shape: {X_test.shape}")
print(f"Training Labels Shape: {y_train.shape}")
print(f"Test Labels Shape: {y_test.shape}")

Training Features Shape: (800, 5)
Test Features Shape: (200, 5)
Training Labels Shape: (800,)
Test Labels Shape: (200,)

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout

# Define the Sequential model
model = Sequential([
    # First hidden layer with 64 neurons and ReLU activation
    Dense(64, activation='relu', input_shape=(X_train.shape[1],)),
    Dropout(0.3), # Dropout with a rate of 30%

    # Second hidden layer with 32 neurons and ReLU activation
    Dense(32, activation='relu'),
    Dropout(0.3), # Dropout with a rate of 30%

    # Output layer for classification
    Dense(len(set(y_encoded)), activation='softmax') # Adjust output
neurons based on the number of classes
])

# Print model summary
model.summary()
```

C:\Users\HP\AppData\Roaming\Python\Python312\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential

```
models, prefer using an `Input(shape)` object as the first layer in the model instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Model: "sequential"
```

Layer (type) Param #	Output Shape
dense (Dense) 384	(None, 64)
dropout (Dropout) 0	(None, 64)
dense_1 (Dense) 2,080	(None, 32)
dropout_1 (Dropout) 0	(None, 32)
dense_2 (Dense) 99	(None, 3)

```
Total params: 2,563 (10.01 KB)
```

```
Trainable params: 2,563 (10.01 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
from tensorflow.keras.optimizers import Adam
```

```
# Compile the model
```

```
model.compile(  
    optimizer=Adam(learning_rate=0.001), # Adam optimizer with a  
    default learning rate of 0.001  
    loss='sparse_categorical_crossentropy', # Loss function for  
    integer-labeled multi-class classification  
    metrics=['accuracy'] # Track accuracy during training  
)
```

```
print("Model compiled successfully!")
```

Model compiled successfully!

```
history = model.fit(X_train, y_train, validation_data=(X_test,  
y_test), epochs=50, batch_size=32)
```

Epoch 1/50

25/25 _____ 1s 18ms/step - accuracy: 0.9283 - loss: 0.1559 - val_accuracy: 0.9700 - val_loss: 0.1012

Epoch 2/50

25/25 _____ 0s 14ms/step - accuracy: 0.9596 - loss: 0.1417 - val_accuracy: 0.9650 - val_loss: 0.0995

Epoch 3/50

25/25 _____ 1s 13ms/step - accuracy: 0.9601 - loss: 0.1201 - val_accuracy: 0.9700 - val_loss: 0.1011

Epoch 4/50

25/25 _____ 0s 13ms/step - accuracy: 0.9377 - loss: 0.1352 - val_accuracy: 0.9700 - val_loss: 0.0975

Epoch 5/50

25/25 _____ 1s 19ms/step - accuracy: 0.9455 - loss: 0.1385 - val_accuracy: 0.9650 - val_loss: 0.0974

Epoch 6/50

25/25 _____ 0s 12ms/step - accuracy: 0.9457 - loss: 0.1391 - val_accuracy: 0.9750 - val_loss: 0.0988

Epoch 7/50

25/25 _____ 0s 14ms/step - accuracy: 0.9503 - loss: 0.1243 - val_accuracy: 0.9750 - val_loss: 0.0952

Epoch 8/50

25/25 _____ 0s 14ms/step - accuracy: 0.9496 - loss: 0.1230 - val_accuracy: 0.9700 - val_loss: 0.0980

Epoch 9/50

25/25 _____ 0s 14ms/step - accuracy: 0.9513 - loss: 0.1445 - val_accuracy: 0.9750 - val_loss: 0.0944

Epoch 10/50

25/25 _____ 0s 13ms/step - accuracy: 0.9529 - loss: 0.1177 - val_accuracy: 0.9750 - val_loss: 0.0948

Epoch 11/50

25/25 _____ 0s 13ms/step - accuracy: 0.9460 - loss: 0.1385 - val_accuracy: 0.9700 - val_loss: 0.0956

Epoch 12/50

25/25 _____ 0s 13ms/step - accuracy: 0.9375 - loss: 0.1427 - val_accuracy: 0.9700 - val_loss: 0.0916

Epoch 13/50

25/25 _____ 0s 14ms/step - accuracy: 0.9740 - loss: 0.0996 - val_accuracy: 0.9800 - val_loss: 0.0926

Epoch 14/50

25/25 _____ 0s 13ms/step - accuracy: 0.9627 - loss: 0.1017 - val_accuracy: 0.9750 - val_loss: 0.0890

Epoch 15/50
25/25 _____ 0s 15ms/step - accuracy: 0.9500 - loss: 0.1165 - val_accuracy: 0.9750 - val_loss: 0.0915
Epoch 16/50
25/25 _____ 0s 14ms/step - accuracy: 0.9543 - loss: 0.1082 - val_accuracy: 0.9800 - val_loss: 0.0927
Epoch 17/50
25/25 _____ 0s 15ms/step - accuracy: 0.9609 - loss: 0.1172 - val_accuracy: 0.9750 - val_loss: 0.0909
Epoch 18/50
25/25 _____ 0s 14ms/step - accuracy: 0.9580 - loss: 0.1167 - val_accuracy: 0.9750 - val_loss: 0.0899
Epoch 19/50
25/25 _____ 0s 14ms/step - accuracy: 0.9501 - loss: 0.1127 - val_accuracy: 0.9700 - val_loss: 0.0960
Epoch 20/50
25/25 _____ 0s 12ms/step - accuracy: 0.9531 - loss: 0.1028 - val_accuracy: 0.9750 - val_loss: 0.0897
Epoch 21/50
25/25 _____ 0s 13ms/step - accuracy: 0.9471 - loss: 0.1192 - val_accuracy: 0.9750 - val_loss: 0.0878
Epoch 22/50
25/25 _____ 0s 14ms/step - accuracy: 0.9601 - loss: 0.1095 - val_accuracy: 0.9750 - val_loss: 0.0907
Epoch 23/50
25/25 _____ 0s 12ms/step - accuracy: 0.9457 - loss: 0.1163 - val_accuracy: 0.9800 - val_loss: 0.0878
Epoch 24/50
25/25 _____ 0s 13ms/step - accuracy: 0.9518 - loss: 0.1308 - val_accuracy: 0.9750 - val_loss: 0.0853
Epoch 25/50
25/25 _____ 0s 13ms/step - accuracy: 0.9637 - loss: 0.1003 - val_accuracy: 0.9750 - val_loss: 0.0848
Epoch 26/50
25/25 _____ 0s 15ms/step - accuracy: 0.9483 - loss: 0.1262 - val_accuracy: 0.9750 - val_loss: 0.0882
Epoch 27/50
25/25 _____ 0s 13ms/step - accuracy: 0.9652 - loss: 0.0844 - val_accuracy: 0.9800 - val_loss: 0.0846
Epoch 28/50
25/25 _____ 0s 15ms/step - accuracy: 0.9600 - loss: 0.1093 - val_accuracy: 0.9750 - val_loss: 0.0850
Epoch 29/50
25/25 _____ 1s 18ms/step - accuracy: 0.9614 - loss: 0.1086 - val_accuracy: 0.9750 - val_loss: 0.0880
Epoch 30/50
25/25 _____ 0s 14ms/step - accuracy: 0.9630 - loss: 0.0922 - val_accuracy: 0.9800 - val_loss: 0.0860
Epoch 31/50

```
25/25 _____ 0s 14ms/step - accuracy: 0.9762 - loss:
0.0741 - val_accuracy: 0.9800 - val_loss: 0.0859
Epoch 32/50
25/25 _____ 0s 14ms/step - accuracy: 0.9690 - loss:
0.0829 - val_accuracy: 0.9700 - val_loss: 0.0864
Epoch 33/50
25/25 _____ 0s 14ms/step - accuracy: 0.9704 - loss:
0.0837 - val_accuracy: 0.9700 - val_loss: 0.0874
Epoch 34/50
25/25 _____ 0s 14ms/step - accuracy: 0.9519 - loss:
0.1070 - val_accuracy: 0.9750 - val_loss: 0.0859
Epoch 35/50
25/25 _____ 0s 14ms/step - accuracy: 0.9699 - loss:
0.0739 - val_accuracy: 0.9750 - val_loss: 0.0938
Epoch 36/50
25/25 _____ 0s 14ms/step - accuracy: 0.9638 - loss:
0.0909 - val_accuracy: 0.9850 - val_loss: 0.0877
Epoch 37/50
25/25 _____ 0s 13ms/step - accuracy: 0.9705 - loss:
0.0832 - val_accuracy: 0.9750 - val_loss: 0.0870
Epoch 38/50
25/25 _____ 0s 13ms/step - accuracy: 0.9747 - loss:
0.0769 - val_accuracy: 0.9750 - val_loss: 0.0869
Epoch 39/50
25/25 _____ 0s 13ms/step - accuracy: 0.9722 - loss:
0.0748 - val_accuracy: 0.9750 - val_loss: 0.0856
Epoch 40/50
25/25 _____ 1s 20ms/step - accuracy: 0.9642 - loss:
0.0912 - val_accuracy: 0.9800 - val_loss: 0.0871
Epoch 41/50
25/25 _____ 0s 13ms/step - accuracy: 0.9680 - loss:
0.0753 - val_accuracy: 0.9750 - val_loss: 0.0875
Epoch 42/50
25/25 _____ 0s 14ms/step - accuracy: 0.9625 - loss:
0.0880 - val_accuracy: 0.9700 - val_loss: 0.0902
Epoch 43/50
25/25 _____ 0s 13ms/step - accuracy: 0.9685 - loss:
0.0803 - val_accuracy: 0.9750 - val_loss: 0.0878
Epoch 44/50
25/25 _____ 0s 15ms/step - accuracy: 0.9798 - loss:
0.0676 - val_accuracy: 0.9700 - val_loss: 0.0921
Epoch 45/50
25/25 _____ 0s 13ms/step - accuracy: 0.9643 - loss:
0.0773 - val_accuracy: 0.9800 - val_loss: 0.0890
Epoch 46/50
25/25 _____ 0s 14ms/step - accuracy: 0.9701 - loss:
0.0820 - val_accuracy: 0.9750 - val_loss: 0.0886
Epoch 47/50
25/25 _____ 1s 15ms/step - accuracy: 0.9750 - loss:
```

```

0.0884 - val_accuracy: 0.9800 - val_loss: 0.0859
Epoch 48/50
25/25 _____ 0s 12ms/step - accuracy: 0.9723 - loss:
0.0828 - val_accuracy: 0.9800 - val_loss: 0.0938
Epoch 49/50
25/25 _____ 0s 15ms/step - accuracy: 0.9581 - loss:
0.0994 - val_accuracy: 0.9700 - val_loss: 0.0971
Epoch 50/50
25/25 _____ 0s 14ms/step - accuracy: 0.9596 - loss:
0.0830 - val_accuracy: 0.9800 - val_loss: 0.0951

# Evaluate the model on the test dataset
test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=1)

# Print the results
print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")

7/7 _____ 0s 15ms/step - accuracy: 0.9756 - loss:
0.1593
Test Loss: 0.0951
Test Accuracy: 0.9800

# Extract training and validation metrics
training_accuracy = history.history['accuracy'][-1]
validation_accuracy = history.history['val_accuracy'][-1]
training_loss = history.history['loss'][-1]
validation_loss = history.history['val_loss'][-1]

print("Training Accuracy:", training_accuracy)
print("Validation Accuracy:", validation_accuracy)
print("Training Loss:", training_loss)
print("Validation Loss:", validation_loss)

Training Accuracy: 0.9662500023841858
Validation Accuracy: 0.9800000190734863
Training Loss: 0.07806619256734848
Validation Loss: 0.09511351585388184

import matplotlib.pyplot as plt

# Plot training and validation accuracy
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()


```

```

# Plot training and validation loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

# Display the plots
plt.tight_layout()
plt.show()

```



```

# Save the entire model to a file
model.save("soil_classification_model.h5")
print("Model saved as 'soil_classification_model.h5'")

WARNING:absl:You are saving your model as an HDF5 file via
`model.save()` or `keras.saving.save_model(model)`. This file format
is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')` or
`keras.saving.save_model(model, 'my_model.keras')`.

Model saved as 'soil_classification_model.h5'

from tensorflow.keras.models import load_model

# Load the model from the file
loaded_model = load_model("soil_classification_model.h5")
print("Model loaded successfully!")

```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

Model loaded successfully!

Evaluate the loaded model on the test set

```
loss, accuracy = loaded_model.evaluate(X_test, y_test)
print(f"Loaded Model Accuracy: {accuracy:.2f}")
```

Make predictions with the loaded model

```
new_soil_data = np.array([[2, 6.5, 15.0, 3.5, 1.2]]) # Example new data
new_soil_data_scaled = scaler.transform(new_soil_data)
predictions = loaded_model.predict(new_soil_data_scaled)
predicted_class = np.argmax(predictions, axis=1)
print("Predicted Soil Type:", predicted_class)
```

7/7 ————— 1s 13ms/step - accuracy: 0.9262 - loss: 0.2066

Loaded Model Accuracy: 0.93

1/1 ————— 0s 259ms/step

Predicted Soil Type: [0]

```
import os
```

Get the file size of the saved model

```
model_file = "soil_classification_model.h5"
```

Check if the file exists

```
if os.path.exists(model_file):
    model_size = os.path.getsize(model_file) / (1024 * 1024) #
    Convert bytes to MB
    print(f"Model size: {model_size:.2f} MB")
```

```
else:
    print(f"Model file '{model_file}' not found. Please save the model first.")
```

Model size: 0.06 MB